


O'REILLY®

AI Engineering

Building Applications
with Foundation Models



Chip Huyen

"This book offers a comprehensive, well-structured guide to the essential aspects of building generative AI systems. A must-read for any professional looking to scale AI across the enterprise."

Vittorio Cretella, former global CIO at P&G and Mars

"Chip Huyen gets generative AI. She is a remarkable teacher and writer whose work has been instrumental in helping teams bring AI into production. Drawing on her deep expertise, *AI Engineering* is a comprehensive and holistic guide to building generative AI applications in production."

Luke Metz, cocreator of ChatGPT, former research manager at OpenAI

AI Engineering

Foundation models have enabled many new AI use cases while lowering the barriers to entry for building AI products. This has transformed AI from an esoteric discipline into a powerful development tool that anyone can use—including those with no prior AI experience.

In this accessible guide, author Chip Huyen discusses AI engineering: the process of building applications with readily available foundation models. AI application developers will discover how to navigate the AI landscape, including models, datasets, evaluation benchmarks, and the seemingly infinite number of application patterns. The book also introduces a practical framework for developing an AI application and efficiently deploying it.

- Understand what AI engineering is and how it differs from traditional machine learning engineering
- Learn the process for developing an AI application, the challenges at each step, and approaches to address them
- Explore various model adaptation techniques, including prompt engineering, RAG, finetuning, agents, and dataset engineering, and understand how and why they work
- Examine the bottlenecks for latency and cost when serving foundation models and learn how to overcome them
- Choose the right model, metrics, data, and developmental patterns for your needs

Chip Huyen works at the intersection of AI, data, and storytelling. Previously, she was with Snorkel AI and NVIDIA, founded an AI infrastructure startup (acquired), and taught machine learning systems design at Stanford. Her book *Designing Machine Learning Systems* (O'Reilly) has been translated into over 10 languages.

DATA

US \$79.99 CAN \$99.99

ISBN: 978-1-098-16630-4



5 7 9 9 9



O'REILLY®

Praise for *AI Engineering*

This book offers a comprehensive, well-structured guide to the essential aspects of building generative AI systems. A must-read for any professional looking to scale AI across the enterprise.

—Vittorio Cretella, former global CIO, P&G and Mars

Chip Huyen gets generative AI. On top of that, she is a remarkable teacher and writer whose work has been instrumental in helping teams bring AI into production.

Drawing on her deep expertise, *AI Engineering* serves as a comprehensive and holistic guide, masterfully detailing everything required to design and deploy generative AI applications in production.

—Luke Metz, cocreator of ChatGPT,
former research manager at OpenAI

Every AI engineer building real-world applications should read this book. It's a vital guide to end-to-end AI system design, from model development and evaluation to large-scale deployment and operation.

—Andrei Lopatenko, Director Search and AI, Neuron7

This book serves as an essential guide for building AI products that can scale.

Unlike other books that focus on tools or current trends that are constantly changing, Chip delivers timeless foundational knowledge. Whether you're a product manager or an engineer, this book effectively bridges the collaboration gap between cross-functional teams, making it a must-read for anyone involved in AI development.

—Aileen Bui, AI Product Operations Manager, Google

This is the definitive segue into AI engineering from one of the greats of ML engineering! Chip has seen through successful projects and careers at every stage of a company and for the first time ever condensed her expertise for new AI Engineers entering the field.

—swyx, *Curator, AI.Engineer*

AI Engineering is a practical guide that provides the most up-to-date information on AI development, making it approachable for novice and expert leaders alike. This book is an essential resource for anyone looking to build robust and scalable AI systems.

—Vicki Reyzelman, *Chief AI Solutions Architect, Mave Sparks*

AI Engineering is a comprehensive guide that serves as an essential reference for both understanding and implementing AI systems in practice.

—Han Lee, *Director—Data Science, Moody's*

AI Engineering is an essential guide for anyone building software with Generative AI! It demystifies the technology, highlights the importance of evaluation, and shares what should be done to achieve quality before starting with costly fine-tuning.

—Rafal Kawala, *Senior AI Engineering Director, 16 years of experience working in a Fortune 500 company*

AI Engineering

Building Applications with Foundation Models

Chip Huyen

O'REILLY®

AI Engineering

by Chip Huyen

Copyright © 2025 Developer Experience Advisory LLC. All rights reserved.

Printed in the United States of America.

Published by O'Reilly Media, Inc., 1005 Gravenstein Highway North, Sebastopol, CA 95472.

O'Reilly books may be purchased for educational, business, or sales promotional use. Online editions are also available for most titles (<http://oreilly.com>). For more information, contact our corporate/institutional sales department: 800-998-9938 or corporate@oreilly.com.

Acquisitions Editor: Nicole Butterfield

Development Editor: Melissa Potter

Production Editor: Beth Kelly

Copyeditor: Liz Wheeler

Proofreader: Piper Editorial Consulting, LLC

Indexer: WordCo Indexing Services, Inc.

Interior Designer: David Futato

Cover Designer: Karen Montgomery

Illustrator: Kate Dullea

December 2024: First Edition

Revision History for the First Edition

2024-12-04: First Release

See <http://oreilly.com/catalog/errata.csp?isbn=9781098166304> for release details.

The O'Reilly logo is a registered trademark of O'Reilly Media, Inc. *AI Engineering*, the cover image, and related trade dress are trademarks of O'Reilly Media, Inc.

The views expressed in this work are those of the author and do not represent the publisher's views. While the publisher and the author have used good faith efforts to ensure that the information and instructions contained in this work are accurate, the publisher and the author disclaim all responsibility for errors or omissions, including without limitation responsibility for damages resulting from the use of or reliance on this work. Use of the information and instructions contained in this work is at your own risk. If any code samples or other technology this work contains or describes is subject to open source licenses or the intellectual property rights of others, it is your responsibility to ensure that your use thereof complies with such licenses and/or rights.

978-1-098-16630-4

[LSI]

Table of Contents

Preface.....	xi
1. Introduction to Building AI Applications with Foundation Models.....	1
The Rise of AI Engineering	2
From Language Models to Large Language Models	2
From Large Language Models to Foundation Models	8
From Foundation Models to AI Engineering	12
Foundation Model Use Cases	16
Coding	20
Image and Video Production	22
Writing	22
Education	24
Conversational Bots	26
Information Aggregation	26
Data Organization	27
Workflow Automation	28
Planning AI Applications	28
Use Case Evaluation	29
Setting Expectations	32
Milestone Planning	33
Maintenance	34
The AI Engineering Stack	35
Three Layers of the AI Stack	37
AI Engineering Versus ML Engineering	39
AI Engineering Versus Full-Stack Engineering	46
Summary	47

2. Understanding Foundation Models.....	49
Training Data	50
Multilingual Models	51
Domain-Specific Models	56
Modeling	58
Model Architecture	58
Model Size	67
Post-Training	78
Supervised Finetuning	80
Preference Finetuning	83
Sampling	88
Sampling Fundamentals	88
Sampling Strategies	90
Test Time Compute	96
Structured Outputs	99
The Probabilistic Nature of AI	105
Summary	111
3. Evaluation Methodology.....	113
Challenges of Evaluating Foundation Models	114
Understanding Language Modeling Metrics	118
Entropy	119
Cross Entropy	120
Bits-per-Character and Bits-per-Byte	121
Perplexity	121
Perplexity Interpretation and Use Cases	122
Exact Evaluation	125
Functional Correctness	126
Similarity Measurements Against Reference Data	127
Introduction to Embedding	134
AI as a Judge	136
Why AI as a Judge?	137
How to Use AI as a Judge	138
Limitations of AI as a Judge	141
What Models Can Act as Judges?	145
Ranking Models with Comparative Evaluation	148
Challenges of Comparative Evaluation	152
The Future of Comparative Evaluation	155
Summary	156

4. Evaluate AI Systems.....	159
Evaluation Criteria	160
Domain-Specific Capability	161
Generation Capability	163
Instruction-Following Capability	172
Cost and Latency	177
Model Selection	179
Model Selection Workflow	179
Model Build Versus Buy	181
Navigate Public Benchmarks	191
Design Your Evaluation Pipeline	200
Step 1. Evaluate All Components in a System	200
Step 2. Create an Evaluation Guideline	202
Step 3. Define Evaluation Methods and Data	204
Summary	208
5. Prompt Engineering.....	211
Introduction to Prompting	212
In-Context Learning: Zero-Shot and Few-Shot	213
System Prompt and User Prompt	215
Context Length and Context Efficiency	218
Prompt Engineering Best Practices	220
Write Clear and Explicit Instructions	220
Provide Sufficient Context	223
Break Complex Tasks into Simpler Subtasks	224
Give the Model Time to Think	227
Iterate on Your Prompts	229
Evaluate Prompt Engineering Tools	230
Organize and Version Prompts	233
Defensive Prompt Engineering	235
Proprietary Prompts and Reverse Prompt Engineering	236
Jailbreaking and Prompt Injection	238
Information Extraction	243
Defenses Against Prompt Attacks	248
Summary	251
6. RAG and Agents.....	253
RAG	253
RAG Architecture	256
Retrieval Algorithms	257
Retrieval Optimization	267

RAG Beyond Texts	273
Agents	275
Agent Overview	276
Tools	278
Planning	281
Agent Failure Modes and Evaluation	298
Memory	300
Summary	305
7. Finetuning.....	307
Finetuning Overview	308
When to Finetune	311
Reasons to Finetune	311
Reasons Not to Finetune	312
Finetuning and RAG	316
Memory Bottlenecks	319
Backpropagation and Trainable Parameters	320
Memory Math	322
Numerical Representations	325
Quantization	328
Finetuning Techniques	332
Parameter-Efficient Finetuning	332
Model Merging and Multi-Task Finetuning	347
Finetuning Tactics	357
Summary	361
8. Dataset Engineering.....	363
Data Curation	365
Data Quality	368
Data Coverage	369
Data Quantity	372
Data Acquisition and Annotation	377
Data Augmentation and Synthesis	380
Why Data Synthesis	381
Traditional Data Synthesis Techniques	383
AI-Powered Data Synthesis	386
Model Distillation	395
Data Processing	396
Inspect Data	397
Deduplicate Data	399
Clean and Filter Data	401

Format Data	401
Summary	403
9. Inference Optimization.....	405
Understanding Inference Optimization	406
Inference Overview	406
Inference Performance Metrics	412
AI Accelerators	419
Inference Optimization	426
Model Optimization	426
Inference Service Optimization	440
Summary	447
10. AI Engineering Architecture and User Feedback.....	449
AI Engineering Architecture	449
Step 1. Enhance Context	450
Step 2. Put in Guardrails	451
Step 3. Add Model Router and Gateway	456
Step 4. Reduce Latency with Caches	460
Step 5. Add Agent Patterns	463
Monitoring and Observability	465
AI Pipeline Orchestration	472
User Feedback	474
Extracting Conversational Feedback	475
Feedback Design	480
Feedback Limitations	490
Summary	492
Epilogue.....	495
Index.....	497

Preface

When ChatGPT came out, like many of my colleagues, I was disoriented. What surprised me wasn't the model's size or capabilities. For over a decade, the AI community has known that scaling up a model improves it. In 2012, the AlexNet authors noted in **their landmark paper** that: "All of our experiments suggest that our results can be improved simply by waiting for faster GPUs and bigger datasets to become available."^{1, 2}

What surprised me was the sheer number of applications this capability boost unlocked. I thought a small increase in model quality metrics might result in a modest increase in applications. Instead, it resulted in an explosion of new possibilities.

Not only have these new AI capabilities increased the demand for AI applications, but they have also lowered the entry barrier for developers. It's become so easy to get started with building AI applications. It's even possible to build an application without writing a single line of code. This shift has transformed AI from a specialized discipline into a powerful development tool everyone can use.

Even though AI adoption today seems new, it's built upon techniques that have been around for a while. Papers about language modeling came out as early as the 1950s. Retrieval-augmented generation (RAG) applications are built upon retrieval technology that has powered search and recommender systems since long before the term RAG was coined. The best practices for deploying traditional machine learning applications—systematic experimentation, rigorous evaluation, relentless optimization for faster and cheaper models—are still the best practices for working with foundation model-based applications.

1 An author of the AlexNet paper, Ilya Sutskever, went on to cofound OpenAI, turning this lesson into reality with GPT models.

2 Even **my small project in 2017**, which used a language model to evaluate translation quality, concluded that we needed "a better language model."

The familiarity and ease of use of many AI engineering techniques can mislead people into thinking there is nothing new to AI engineering. But while many principles for building AI applications remain the same, the scale and improved capabilities of AI models introduce opportunities and challenges that require new solutions.

This book covers the end-to-end process of adapting foundation models to solve real-world problems, encompassing tried-and-true techniques from other engineering fields and techniques emerging with foundation models.

I set out to write the book because I wanted to learn, and I did learn a lot. I learned from the projects I worked on, the papers I read, and the people I interviewed. During the process of writing this book, I used notes from over 100 conversations and interviews, including researchers from major AI labs (OpenAI, Google, Anthropic, ...), framework developers (NVIDIA, Meta, Hugging Face, Anyscale, LangChain, LlamaIndex, ...), executives and heads of AI/data at companies of different sizes, product managers, community researchers, and independent application developers (see “[Acknowledgments](#)” on page xx).

I especially learned from early readers who tested my assumptions, introduced me to different perspectives, and exposed me to new problems and approaches. Some sections of the book have also received thousands of comments from the community after being shared on [my blog](#), many giving me new perspectives or confirming a hypothesis.

I hope that this learning process will continue for me now that the book is in your hands, as you have experiences and perspectives that are unique to you. Please feel free to share any feedback you might have for this book with me via [X](#), [LinkedIn](#), or email at hi@huyenchip.com.

What This Book Is About

This book provides a framework for adapting foundation models, which include both large language models (LLMs) and large multimodal models (LMMs), to specific applications.

There are many different ways to build an application. This book outlines various solutions and also raises questions you can ask to evaluate the best solution for your needs. Some of the many questions that this book can help you answer are:

- Should I build this AI application?
- How do I evaluate my application? Can I use AI to evaluate AI outputs?
- What causes hallucinations? How do I detect and mitigate hallucinations?
- What are the best practices for prompt engineering?
- Why does RAG work? What are the strategies for doing RAG?

- What's an agent? How do I build and evaluate an agent?
- When to finetune a model? When not to finetune a model?
- How much data do I need? How do I validate the quality of my data?
- How do I make my model faster, cheaper, and secure?
- How do I create a feedback loop to improve my application continually?

The book will also help you navigate the overwhelming AI landscape: types of models, evaluation benchmarks, and a seemingly infinite number of use cases and application patterns.

The content in this book is illustrated using case studies, many of which I worked on, backed by ample references and extensively reviewed by experts from a wide range of backgrounds. Although the book took two years to write, it draws from my experience working with language models and ML systems from the last decade.

Like my previous O'Reilly book, *Designing Machine Learning Systems* (DMLS), this book focuses on the fundamentals of AI engineering instead of any specific tool or API. Tools become outdated quickly, but fundamentals should last longer.³

Reading *AI Engineering* (AIE) with *Designing Machine Learning Systems* (DMLS)

AIE can be a companion to DMLS. DMLS focuses on building applications on top of traditional ML models, which involves more tabular data annotations, feature engineering, and model training. AIE focuses on building applications on top of foundation models, which involves more prompt engineering, context construction, and parameter-efficient finetuning. Both books are self-contained and modular, so you can read either book independently.

Since foundation models are ML models, some concepts are relevant to working with both. If a topic is relevant to AIE but has been discussed extensively in DMLS, it'll still be covered in this book, but to a lesser extent, with pointers to relevant resources.

Note that many topics are covered in DMLS but not in AIE, and vice versa. The first chapter of this book also covers the differences between traditional ML engineering and AI engineering. A real-world system often involves both traditional ML models and foundation models, so knowledge about working with both is often necessary.

³ Teaching a course on how to use TensorFlow in 2017 taught me a painful lesson about how quickly tools and tutorials become outdated.

Determining whether something will last, however, is often challenging. I relied on three criteria. First, for a problem, I determined whether it results from the fundamental limitations of how AI works or if it'll go away with better models. If a problem is fundamental, I'll analyze its challenges and solutions to address each challenge. I'm a fan of the start-simple approach, so for many problems, I'll start from the simplest solution and then progress with more complex solutions to address rising challenges.

Second, I consulted an extensive network of researchers and engineers, who are smarter than I am, about what they think are the most important problems and solutions.

Occasionally, I also relied on **Lindy's Law**, which infers that the future life expectancy of a technology is proportional to its current age. So if something has been around for a while, I assume that it'll continue existing for a while longer.

In this book, however, I occasionally included a concept that I believe to be temporary because it's immediately useful for some application developers or because it illustrates an interesting problem-solving approach.

What This Book Is Not

This book isn't a tutorial. While it mentions specific tools and includes pseudocode snippets to illustrate certain concepts, it doesn't teach you how to use a tool. Instead, it offers a framework for selecting tools. It includes many discussions on the trade-offs between different solutions and the questions you should ask when evaluating a solution. When you want to use a tool, it's usually easy to find tutorials for it online. AI chatbots are also pretty good at helping you get started with popular tools.

This book isn't an ML theory book. It doesn't explain what a neural network is or how to build and train a model from scratch. While it explains many theoretical concepts immediately relevant to the discussion, the book is a practical book that focuses on helping you build successful AI applications to solve real-world problems.

While it's possible to build foundation model-based applications without ML expertise, a basic understanding of ML and statistics can help you build better applications and save you from unnecessary suffering. You can read this book without any prior ML background. However, you will be more effective while building AI applications if you know the following concepts:

- Probabilistic concepts such as sampling, determinism, and distribution.
- ML concepts such as supervision, self-supervision, log-likelihood, gradient descent, backpropagation, loss function, and hyperparameter tuning.

- Various neural network architectures, including feedforward, recurrent, and transformer.
- Metrics such as accuracy, F1, precision, recall, cosine similarity, and cross entropy.

If you don't know them yet, don't worry—this book has either brief, high-level explanations or pointers to resources that can get you up to speed.

Who This Book Is For

This book is for anyone who wants to leverage foundation models to solve real-world problems. This is a technical book, so the language of this book is geared toward technical roles, including AI engineers, ML engineers, data scientists, engineering managers, and technical product managers. This book is for you if you can relate to one of the following scenarios:

- You're building or optimizing an AI application, whether you're starting from scratch or looking to move beyond the demo phase into a production-ready stage. You may also be facing issues like hallucinations, security, latency, or costs, and need targeted solutions.
- You want to streamline your team's AI development process, making it more systematic, faster, and reliable.
- You want to understand how your organization can leverage foundation models to improve the business's bottom line and how to build a team to do so.

You can also benefit from the book if you belong to one of the following groups:

- Tool developers who want to identify underserved areas in AI engineering to position your products in the ecosystem.
- Researchers who want to better understand AI use cases.
- Job candidates seeking clarity on the skills needed to pursue a career as an AI engineer.
- Anyone wanting to better understand AI's capabilities and limitations, and how it might affect different roles.

I love getting to the bottom of things, so some sections dive a bit deeper into the technical side. While many early readers like the detail, it might not be for everyone. I'll give you a heads-up before things get too technical. Feel free to skip ahead if it feels a little too in the weeds!

Navigating This Book

This book is structured to follow the typical process for developing an AI application. Here's what this typical process looks like and how each chapter fits into the process. Because this book is modular, you're welcome to skip any section that you're already familiar with or that is less relevant to you.

Before deciding to build an AI application, it's necessary to understand what this process involves and answer questions such as: Is this application necessary? Is AI needed? Do I have to build this application myself? The first chapter of the book helps you answer these questions. It also covers a range of successful use cases to give a sense of what foundation models can do.

While an ML background is not necessary to build AI applications, understanding how a foundation model works under the hood is useful to make the most out of it. **Chapter 2** analyzes the making of a foundation model and the design decisions with significant impacts on downstream applications, including its training data recipe, model architectures and scales, and how the model is trained to align to human preference. It then discusses how a model generates a response, which helps explain the model's seemingly baffling behaviors, like inconsistency and hallucinations. Changing the generation setting of a model is also often a cheap and easy way to significantly boost the model's performance.

Once you've committed to building an application with foundation models, evaluation will be an integral part of every step along the way. Evaluation is one of the hardest, if not the hardest, challenges of AI engineering. This book dedicates two chapters, **Chapters 3 and 4**, to explore different evaluation methods and how to use them to create a reliable and systematic evaluation pipeline for your application.

Given a query, the quality of a model's response depends on the following aspects (outside of the model's generation setting):

- The instructions for how the model should behave
- The context the model can use to respond to the query
- The model itself

The next three chapters of the book focus on how to optimize each of these aspects to improve a model's performance for an application. **Chapter 5** covers prompt engineering, starting with what a prompt is, why prompt engineering works, and prompt engineering best practices. It then discusses how bad actors can exploit your application with prompt attacks and how to defend your application against them.

Chapter 6 explores why context is important for a model to generate accurate responses. It zooms into two major application patterns for context construction: RAG and agentic. The RAG pattern is better understood and has proven to work well in

production. On the other hand, while the agentic pattern promises to be much more powerful, it's also more complex and is still being explored.

Chapter 7 is about how to adapt a model to an application by changing the model itself with finetuning. Due to the scale of foundation models, native model finetuning is memory-intensive, and many techniques are developed to allow finetuning better models with less memory. The chapter covers different finetuning approaches, supplemented by a more experimental approach: model merging. This chapter contains a more technical section that shows how to calculate the memory footprint of a model.

Due to the availability of many finetuning frameworks, the finetuning process itself is often straightforward. However, getting data for finetuning is hard. The next chapter is all about data, including data acquisition, data annotations, data synthesis, and data processing. Many of the topics discussed in **Chapter 8** are relevant beyond finetuning, including the question of what data quality means and how to evaluate the quality of your data.

If Chapters 5 to 8 are about improving a model's quality, **Chapter 9** is about making its inference cheaper and faster. It discusses optimization both at the model level and inference service level. If you're using a model API—i.e., someone else hosts your model for you—this API will likely take care of inference optimization for you. However, if you host the model yourself—either an open source model or a model developed in-house—you'll need to implement many of the techniques discussed in this chapter.

The last chapter in the book brings together the different concepts from this book to build an application end-to-end. The second part of the chapter is more product-focused, with discussions on how to design a user feedback system that helps you collect useful feedback while maintaining a good user experience.



I often use “we” in this book to mean you (the reader) and I. It's a habit I got from my teaching days, as I saw writing as a shared learning experience for both the writer and the readers.

Conventions Used in This Book

The following typographical conventions are used in this book:

Italic

Indicates new terms, URLs, email addresses, filenames, and file extensions.

Constant width

Used for program listings, as well as within paragraphs to refer to program elements such as variable or function names, databases, data types, environment variables, statements, input prompts into models, and keywords.

Constant width bold

Shows commands or other text that should be typed literally by the user.

Constant width italic

Shows text that should be replaced with user-supplied values or by values determined by context.



This element signifies a tip or suggestion.



This element signifies a general note.



This element indicates a warning or caution.

Using Code Examples

Supplemental material (code examples, exercises, etc.) is available for download at <https://github.com/chiphuyen/aie-book>. The repository contains additional resources about AI engineering, including important papers and helpful tools. It also covers topics that are too deep to go into in this book. For those interested in the process of writing this book, the GitHub repository also contains behind-the-scenes information and statistics about the book.

If you have a technical question or a problem using the code examples, please send email to support@oreilly.com.

This book is here to help you get your job done. In general, if example code is offered with this book, you may use it in your programs and documentation. You do not need to contact us for permission unless you're reproducing a significant portion of the code. For example, writing a program that uses several chunks of code from